

1.5 Hands-on Lab: Set up a GPU-powered VM on AWS with NVIDIA drivers

What you'll do

Launch an EC2 GPU instance, SSH in, verify the GPU/driver, install Docker + NVIDIA Container Toolkit, and run a quick GPU test in a container.

Prerequisites (5–10 min)

- An AWS account with permissions to launch EC2.
- A key pair (PEM) and a default VPC/subnet.
- Basic terminal/SSH access.

Tip: If this is your first time using GPU instances, you may need a service quota increase before you can launch p4/p5/g6. AWS support can raise it.

Instance/GPU mapping (for choosing later):

- **g6** → NVIDIA **L4** (great for inference/video) ([Amazon Web Services, Inc.](#))
- **g5** → NVIDIA **A10G** (graphics/inference) ([Amazon Web Services, Inc., NVIDIA Developer](#))
- **p4d** → NVIDIA **A100** (training) ([Amazon Web Services, Inc., NVIDIA Blog](#))
- **p5/p5e/p5en** → NVIDIA **H100/H200** (frontier training) ([Amazon Web Services, Inc.](#))

Part A – Launch the EC2 GPU instance (Console) (8–12 min)

1. Pick a Region

Choose an AWS region where your target instance family is available (e.g., us-east-1).

2. AMI (the easy path):

In EC2 → *Launch instance* → **Application and OS Images (Amazon Machine Image)**, search for:

- **“Deep Learning Base GPU AMI (Amazon Linux 2/2023)”** or a framework AMI (e.g., “Deep Learning AMI GPU PyTorch”). These AMIs include NVIDIA drivers/CUDA and are maintained by AWS. ([AWS Documentation](#))

Prefer this AMI for a smoother first run. (You can always use vanilla Ubuntu later and install drivers manually.)

3. Instance type

Pick one that matches your workload and budget, e.g., **g6.xlarge** (L4) for lightweight GPU work, **p4d.24xlarge** (A100) or **p5.48xlarge** (H100) for heavy training. ([Vantage Instances, Amazon Web Services, Inc.](#))

4. Key pair & network

Select your key pair. Use the default VPC/subnet. Create a **Security Group** that allows **SSH (TCP 22)** from your IP.

5. Storage

Set **50–200 GB** gp3 EBS (containers + datasets need space).

6. Launch

Click **Launch instance** and wait until the state is **Running**.

Part B – Connect & verify the GPU (3–5 min)

1. SSH in

```
ssh -i /path/to/your-key.pem ec2-user@<EC2_PUBLIC_IP> # Amazon Linux  
# or  
ssh -i /path/to/your-key.pem ubuntu@<EC2_PUBLIC_IP> # Ubuntu-based AMIs
```

2. Check the GPU & driver

```
nvidia-smi
```

You should see the GPU model (L4/A10G/A100/H100), driver version, and CUDA version. If you used a Deep Learning Base GPU AMI, drivers are already present. ([AWS Documentation](#))

If `nvidia-smi` is missing on a non-DLAMI image, install the correct driver/CUDA using NVIDIA's Linux guide (see "Optional: manual driver install" below). ([NVIDIA Docs](#))

Part C – Install Docker & NVIDIA Container Toolkit (8–12 min)

Skip Docker if your AMI already includes it—otherwise:

Amazon Linux 2023 / 2

```
sudo yum update -y  
sudo yum install -y docker  
sudo systemctl enable --now docker  
sudo usermod -aG docker $USER  
newgrp docker
```

Ubuntu

```
sudo apt-get update  
sudo apt-get install -y docker.io  
sudo usermod -aG docker $USER  
newgrp docker
```

Add the NVIDIA Container Toolkit (official method)

```
# 1) Add repo + key (Ubuntu/Debian example)  
curl -fsSL https://nvidia.github.io/libnvidia-container/gpgkey | \  
  sudo gpg --dearmor -o /usr/share/keyrings/nvidia-container-toolkit-keyring.gpg  
curl -s -L https://nvidia.github.io/libnvidia-container/stable/deb/nvidia-container-toolkit.list | \  
  sed 's#deb https://#deb [signed-by=/usr/share/keyrings/nvidia-container-toolkit-keyring.gpg] https://#g' | \  
  sudo tee /etc/apt/sources.list.d/nvidia-container-toolkit.list  
  
sudo apt-get update  
sudo apt-get install -y nvidia-container-toolkit  
  
# 2) Configure Docker to use NVIDIA runtime  
sudo nvidia-ctk runtime configure --runtime=docker  
sudo systemctl restart docker
```

Part D – Quick GPU test inside a container (2–4 min)

Run a CUDA container and check `nvidia-smi` **from inside** the container:

```
docker run --rm --gpus all nvidia/cuda:12.5.0-runtime-ubuntu22.04 nvidia-smi
```

You should see the same GPU table (driver/library versions) printed by the container—confirming Docker + GPU pass-through works. (If you prefer NGC images, use `nvcr.io/...` after logging in with an NGC API key.) ([NVIDIA Docs](#))

Optional: Manual driver/CUDA install (if you didn't use a DLAMI)

If you launched a plain Ubuntu/AL2023 image and need to install drivers/CUDA yourself, follow NVIDIA's **CUDA Installation Guide for Linux** to match your driver to the CUDA/toolkit version, then reboot and re-run `nvidia-smi`. ([NVIDIA Docs](#))

Troubleshooting (quick hits)

- **`nvidia-smi` not found** → Driver isn't installed or path isn't set. Install via DLAMI or NVIDIA's Linux guide, then reboot. ([AWS Documentation](#), [NVIDIA Docs](#))
 - **Container can't see GPUs** → Make sure **NVIDIA Container Toolkit** is installed and `nvidia-ctk runtime configure --runtime=docker` ran; restart Docker. ([NVIDIA Docs](#))
 - **Wrong instance for your workload** → Use **g6 (L4)** for inference/video, **p4d (A100)** or **p5/p5e (H100/H200)** for training. ([Amazon Web Services, Inc.](#))
-

Clean up (important)

When you're done, **Stop** the instance to pause charges (EBS still billed) or **Terminate** it to delete compute + storage and stop all charges.

You're done 

You now have an AWS GPU VM with drivers and containerized GPU access—ready to pull Triton, PyTorch/TensorFlow, DeepStream, or your own images and start building!